

# **Statistical Data Mining**

## **Spring**

### **Assignment -1**

**Name : Kartik Bapna**

**UB ID :50291058**

**Class No 05**

The State University of New York at Buffalo

Engineering Sciences - Data Science

**Question 1)** Consider the MovieLens data in the “recommenderlab” package.  
Design and evaluate your own recommendation system based

For each user “i” and each movie “j” they did not see, find top “k” most similar users to “i” who have seen “j” and then use them to infer the user “i” ’s rating on movie. Handle all exceptions in a reasonable way and report your strategy if you did so; e.g., if you cannot find “k” users for some movie “j”, then take all users who have seen it.

### Solution 1

Step1 : Uploading and Analyzing data

```
movie_data <- MovieLense

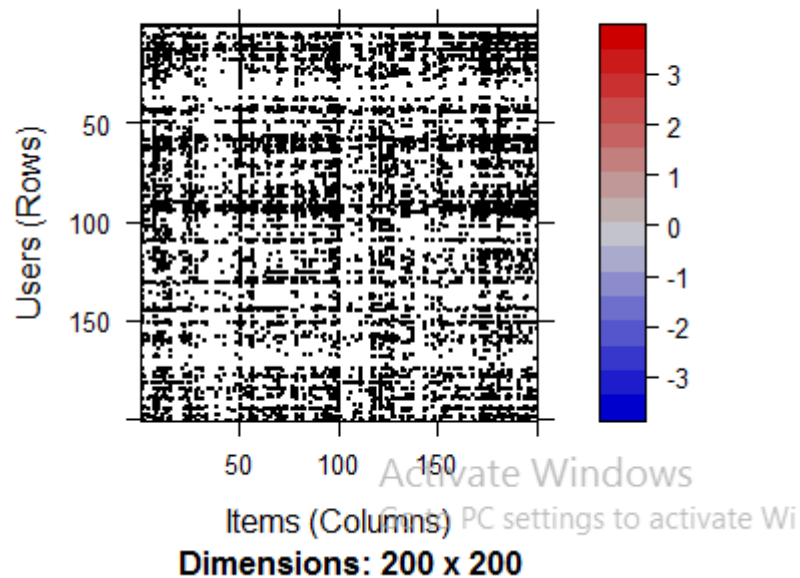
## analyzing data
head(movie_data)

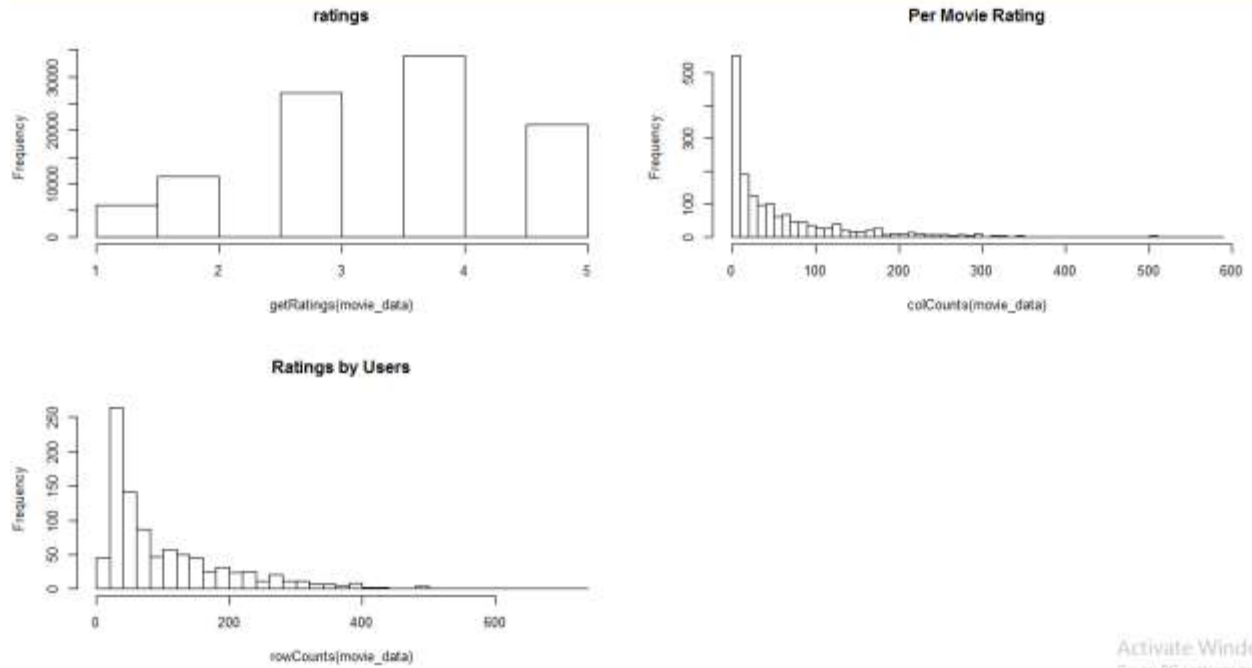
dim(movie_data) ## 943* 1664

image(movie_data[1:200,1:200], main= "Intial Ratings")
```

Step 2 : Scaling the data and Making Plots and visualization data

### Normalized Scaled version of the Ratings





### Step 3 Making Recommender using UBCF data

```
recomm_movie <- Recommender(movie_data, method= "UBCF")
summary(recomm_movie)

> summary(recomm_movie)
      Length      Class      Mode 
      1 Recommender      S4
```

### Step 4 Predicting top 7 movies

```
# predicting best 7 values

best_pred <- predict(recomm_movie, movie_data, n = 7)
best_movie <- bestN(best_pred, n=7)
best_movies_pred <- as(best_movie, "list")
best_movies_pred[1:7]
```

```

$`1`
[1] "Titanic (1997)"          "Air Force One (1997)"    "English Patient, The (1996)"
[4] "Game, The (1997)"       "Rainmaker, The (1997)"  "Wedding Singer, The (1998)"
[7] "Anna Karenina (1997)"

$`2`
[1] "Return of the Jedi (1983)"
[2] "Graduate, The (1967)"
[3] "Blade Runner (1982)"
[4] "Schindler's List (1993)"
[5] "Casablanca (1942)"
[6] "Dr. Strangelove or: How I Learned to Stop Worrying and Love the Bomb (1963)"
[7] "Silence of the Lambs, The (1991)"

$`3`
[1] "Raiders of the Lost Ark (1981)"
[2] "Blade Runner (1982)"
[3] "Star Wars (1977)"
[4] "Empire Strikes Back, The (1980)"
[5] "Dr. Strangelove or: How I Learned to Stop Worrying and Love the Bomb (1963)"
[6] "Wrong Trousers, The (1993)"
[7] "Close Shave, A (1995)"

```

## Part B

**Test the performance of your system using cross-validation. For each data set, the MovieLens database already provides a split of the initial data set into  $N = 5$  folds. This means you will run your algorithm  $N$  times; in each step, use the training partition to make predictions for each user on all terms rated in the test partition (by that user). When you complete all  $N$  iterations, you will have a large number of user-movie pairs from the 5 test partitions on which you can evaluate the performance of your system. Measure the performance of your recommendation system**

Step 5 Apply cross validation

```

# cross validation |
eval_cal<- evaluationScheme(movie_data,method = "cross-validation",given = 15, train=0.5, goodRating=4, k=5)

# recommender model

model<- Recommender(getData(eval_cal,"train"), "UBCF")
summary(model)

```

Step 6 Making model and see prediction

```
model<- Recommender(getData(eval_cal,"train"), "UBCF")
summary(model)

Prediction<- predict(model, getData(eval_cal, "known"), type="ratings")
|
# converting in matrix
new_pred <- as(Prediction, "matrix")
```

#### Step 7 Error Calculation

```
#Top 5|
new_pred[1:5,1:5]

error_cal<- rbind(UBCF = calcPredictionAccuracy(Prediction, getData(eval_cal,"unknown")))
error_cal
```

Measuring the performance of RMSE is 1.09

```
> error_cal
      RMSE      MSE      MAE
[1,] 1.090735 1.189702 0.874382
> |
```

**Question 2) Consider the following ratings table between five users and six items.**

Step 1) Make an excel file and read the file in R

	User	V1	V2	V3	V4	V5	V6
1	1	5	6	7	4	3	NA
2	2	4	NA	3	NA	5	4
3	3	NA	3	4	1	1	NA
4	4	7	4	3	6	NA	4
5	5	1	NA	3	2	2	5

Step 2) Convert the file into matrix form

Step3 )Convert to rating matrix

```
# convert to rating matrix
rrm_matrix <- as(matrix,"realRatingMatrix")
```

Step 4 )Predict the value using pearson method

```
pearson_recc <- Recommender(rrm_matrix, method = "UBCF",param= list(method= "Pearson"))
pearson_predict <- predict(pearson_recc, rrm_matrix, type="ratings")
```

step 5 ) Check the matrix with predicted value for user 2 and others

as we can see all the missing value are getting predicted

	User	V1	V2	V3	V4	V5	V6
[1,]	NA	NA	NA	NA	NA	NA	4.611513
[2,]	NA	NA	3.842212	NA	3.446494	NA	NA
[3,]	NA	2.499412	NA	NA	NA	NA	2.793285
[4,]	NA	NA	NA	NA	NA	4.37253	NA
[5,]	NA	NA	3.290321	NA	NA	NA	NA

Step 6) Repeat the same process for cosine and predict the values for user 2 and others

```
### predicting using cosine
cos_recc <- Recommender(rrm_matrix, method = "IBCF",param= list(method= "Cosine"))
cosine_predict <- predict(cos_recc, rrm_matrix, type="ratings")
as(cosine_predict,"matrix")
```

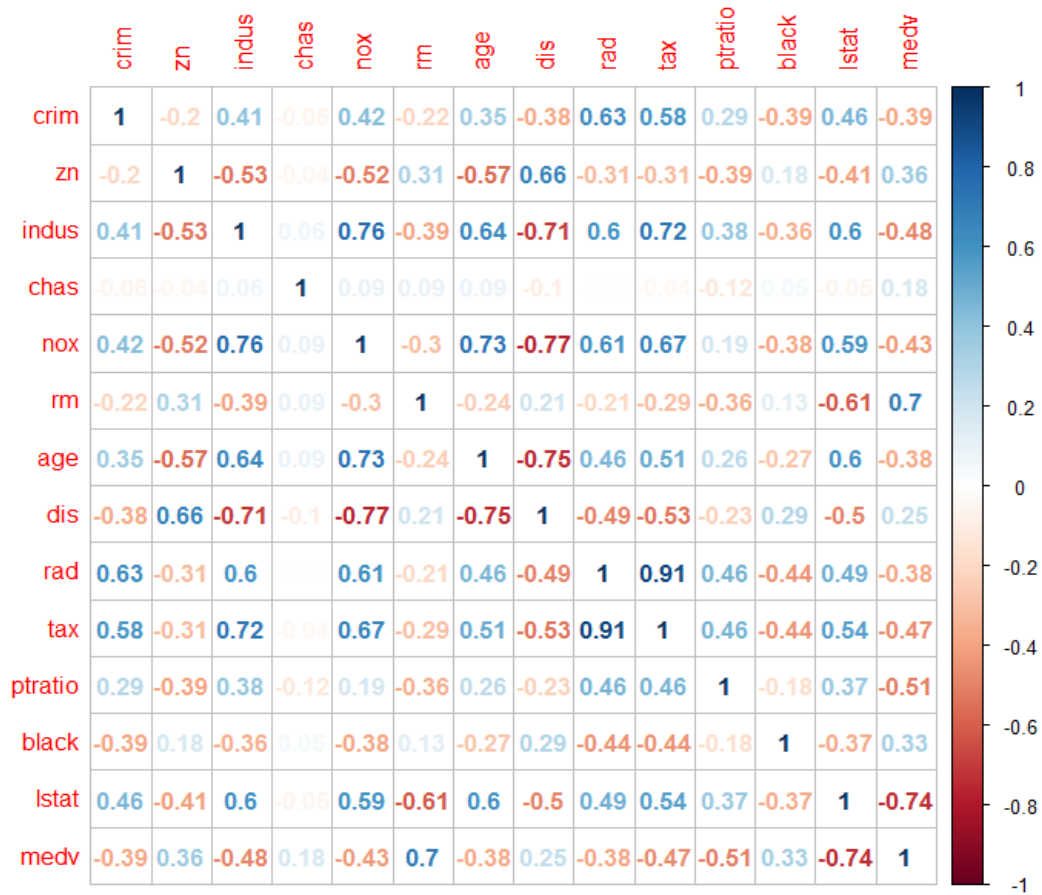
```
as(cosine_predict,"matrix")
```

	User	V1	V2	V3	V4	V5	V6
[1,]	NA	NA	NA	NA	NA	NA	4.273599
[2,]	NA	NA	3.584778	NA	3.891537	NA	NA
[3,]	NA	1.96269	NA	NA	NA	NA	2.286019
[4,]	NA	NA	NA	NA	NA	4.598693	NA
[5,]	NA	NA	3.521462	NA	NA	NA	NA

**question (3) (10 points)** Consider the Boston Housing Data. This data can be accessed in the ElemStatLearn package (available through CRAN).

Solution 3)


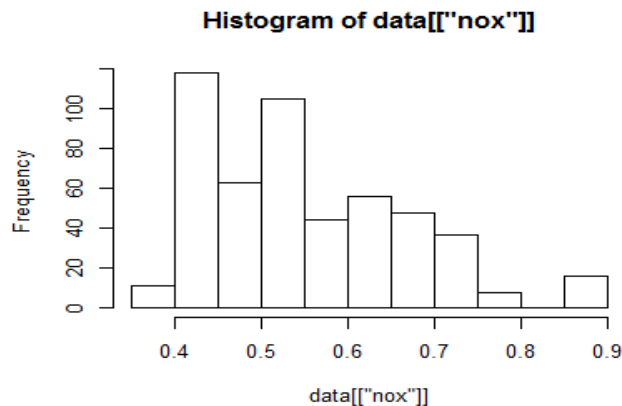
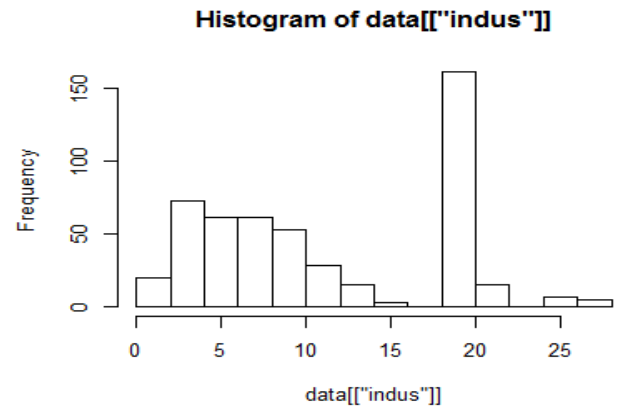

Step1 : Load the Boston data and see the correlation plot



Step 2: Removing the chas as it is not correlated with respect to other variable and plotting histogram for other variables

### Part a

Visualize the data using histograms of the different variables in the data set. Transform the data into a binary incidence matrix, and justify the choices you make in grouping categories.




As per the above histogram the categories are classified as Low, Mid and High as we can see for all the four histogram there are values in low level, mid level and then High Level


Name	Categorised
crim	Low, Mid, High
Indus	Low, Mid, High
Nox	Low, Mid, High
Rm	Low, Mid, High




Histogram of data[["age"]]




Histogram of data[["dis"]]



Histogram of data[["rad"]]




Histogram of data[["tax"]]




As per the Histogram this is classified as , However for Dis attribute it is converted into 4 categories as the value are fluctuating within the histogram and for rad it is converted in two level only

Name	Categorised
age	Low, Mid, High
dis	very low, Low, Mid, High
rad	near , extreme
tax	Low, Mid, High


Histogram of data[["ptratio"]]




Histogram of data[["black"]]



Histogram of data[["lstat"]]



Histogram of data[["medv"]]



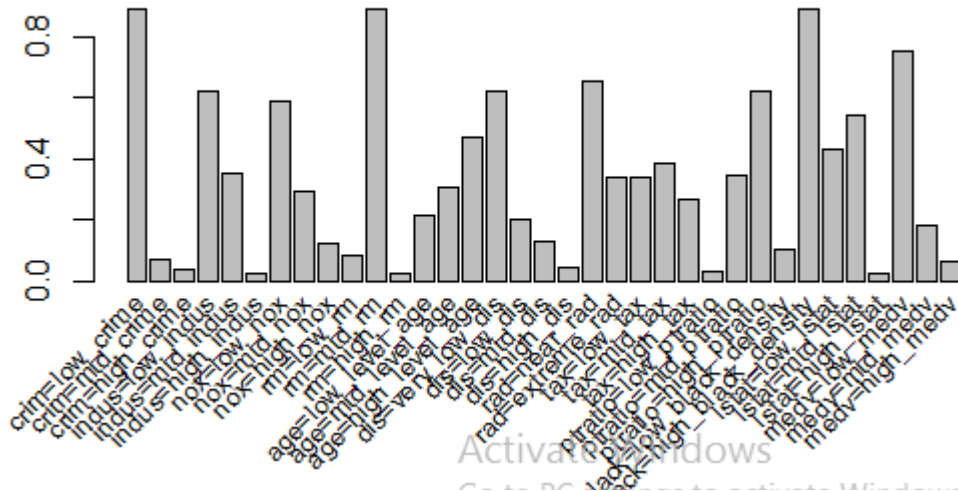
Based on the Histogram the attributes are classified in low, mid and high level however for black and lstat we can categorize them in two level

Name	Categorised
ptratio	Low, Mid, High
black	Low, High
lstat	near , extreme
Medv	Low, Mid, High

Step 3 : Converting to binary incidence matrix

**Part b**

Visualize the data using the itemFrequencyPlot in the “arules” package.  
Apply the apriori algorithm (Do not forget to specify parameters in your write up).

**Step 4 : apply apriori rules**

```
> rules<-apriori(boston_mat,parameter = list(support=.005,confidence =.5))
Apriori

Parameter specification:
 confidence minval  smax  arem  aval originalsupport  maxtime support  minlen maxlen target  ext
      0.5      0.1    1 none FALSE               TRUE         5   0.005      1    10 rules FALSE

Algorithmic control:
 filter tree heap memopt load sort verbose
  0.1 TRUE TRUE  FALSE TRUE    2    TRUE

Absolute minimum support count: 2

set item appearances ...[0 item(s)] done [0.00s].
set transactions ...[35 item(s), 506 transaction(s)] done [0.00s].
sorting and recoding items ... [35 item(s)] done [0.00s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 3 4 5 6 7 8 9 10 done [0.08s].
writing ... [693893 rule(s)] done [0.32s].
creating S4 object ... done [0.50s].
```

Step 5: Checking summary of Rules

set of 693893 rules

```
rule length distribution (lhs + rhs):sizes
  1      2      3      4      5      6      7      8      9      10
10     344    4004   22070   69325  136906  178584  156666   91653  34331

  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  1.00   6.00   7.00   7.12   8.00   10.00
```

```
summary of quality measures:
  support      confidence      lift      count
Min.   :0.005929  Min.   :0.5000  Min.   : 0.5597  Min.   : 3.00
1st Qu.:0.007905  1st Qu.:0.8000  1st Qu.: 1.1195  1st Qu.: 4.00
Median :0.013834  Median :1.0000  Median : 1.6013  Median : 7.00
Mean   :0.026168  Mean   :0.8923  Mean   : 2.0616  Mean   :13.24
3rd Qu.:0.027668  3rd Qu.:1.0000  3rd Qu.: 2.1013  3rd Qu.:14.00
Max.   :0.893281  Max.   :1.0000  Max.   :42.1667  Max.   :452.00
```

```
mining info:
  data ntransactions support confidence
boston_mat      506    0.005      0.5
```

### Part c

A student is interested in a low crime area as close to the city as possible (as measured by "dis"). What can you advise on this matter through the mining of association rules?

```
> rules_1 <- subset(rules, subset = rhs %in% "crim=low_crime" & lhs %in% "dis=very_low_dis" & lift > .5)
> summary(rules_1)
set of 19090 rules

rule length distribution (lhs + rhs):sizes
  2      3      4      5      6      7      8      9      10
1    27    258   1210   3139   4940   4964   3227   1324

  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 2.000   7.000   7.000   7.469   8.000  10.000

summary of quality measures:
  support      confidence      lift      count
Min.   :0.005929  Min.   :0.5000  Min.   :0.5597  Min.   : 3.00
1st Qu.:0.007905  1st Qu.:1.0000  1st Qu.:1.1195  1st Qu.: 4.00
Median :0.013834  Median :1.0000  Median :1.1195  Median : 7.00
Mean   :0.024446  Mean   :0.9623  Mean   :1.0773  Mean   :12.37
3rd Qu.:0.027668  3rd Qu.:1.0000  3rd Qu.:1.1195  3rd Qu.:14.00
Max.   :0.517787  Max.   :1.0000  Max.   :1.1195  Max.   :262.00

mining info:
  data ntransactions support confidence
boston_mat      506    0.005      0.5
```

```
> rules_1 <- subset(rules, subset = rhs %in% "crim=low_crime" & lhs %in% "dis=very_low_dis" & lift >.5)
> inspect(head(sort(rules_1, by = 'lift'), n = 6))
```

	lhs	rhs	support	confidence
[1]	{indus=high_indus,dis=very_low_dis}	=> {crim=low_crime}	0.02371542	1
[2]	{rm=high_rm,dis=very_low_dis}	=> {crim=low_crime}	0.02173913	1
[3]	{dis=very_low_dis,ptratio=low_ptratio}	=> {crim=low_crime}	0.02371542	1
[4]	{dis=very_low_dis,medv=high_medv}	=> {crim=low_crime}	0.04347826	1
[5]	{age=low_level_age,dis=very_low_dis}	=> {crim=low_crime}	0.02371542	1
[6]	{dis=very_low_dis,tax=low_tax}	=> {crim=low_crime}	0.14426877	1

	lift	count
[1]	1.119469	12
[2]	1.119469	11
[3]	1.119469	12
[4]	1.119469	22
[5]	1.119469	12
[6]	1.119469	73

From the above we can see the for low crime area as close to the city(very\_low\_dist) the top rule suggest as per

rule 1) Indus needs to be high as per that select area where proportion of non-retail business acres per town as Indus is high

rule2 ) Ptratio should be low select area where pupil-teacher ratio by town which is low

rule 3) medv should be high select an area where Median value of owner-occupied homes in \$1000's is high and same can inferred from other rules

## part d

A family is moving to the area, and has made schooling a priority. They want schools with low pupil-teacher ratios. What can you advise on this matter through the mining of association rules ?

```
> #part d
> rulesLowPTRatio <- subset(rules, subset = rhs %in% "ptratio=low_ptratio" & lift >.7)
> inspect(head(sort(rulesLowPTRatio, by = 'lift'), n = 6))
```

	lhs	rhs	support	confidence	lift	count
[1]	{nox=mid_nox, rm=high_rm, tax=low_tax}	=> {ptratio=low_ptratio}	0.005928854	1	33.73333	3
[2]	{indus=low_indus, nox=mid_nox, rm=high_rm}	=> {ptratio=low_ptratio}	0.005928854	1	33.73333	3
[3]	{nox=mid_nox, tax=low_tax, medv=high_medv}	=> {ptratio=low_ptratio}	0.009881423	1	33.73333	5
[4]	{indus=low_indus, nox=mid_nox, medv=high_medv}	=> {ptratio=low_ptratio}	0.009881423	1	33.73333	5
[5]	{age=high_level_age, tax=low_tax, medv=high_medv}	=> {ptratio=low_ptratio}	0.005928854	1	33.73333	3

**As per rule 1** we can say the family should select area where nitric oxides concentration is mid level ,average number of dwelling (rm) is high and full-value property-tax rate per \$10,000 is low

**As per rule 2** we can say where proportion of non-retail business acres per town (indus) is low , nitric oxides concentration is mid level and average number of dwelling (rm) is high

Similar inferences can be seen from other rules

### Question e

Use a regression model to solve part d. Are your results comparable? Which provides an easier interpretation? When would regression be preferred, and when would association models be preferred?

After applying regression

```
Call:
lm(formula = Boston$ptratio ~ ., data = Boston)

Residuals:
    Min       1Q   Median       3Q      Max
-4.1190 -1.0126 -0.0060  0.8961  4.8945

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.484e+01  1.352e+00  18.379 < 2e-16 ***
crim         -1.578e-02  1.085e-02  -1.454  0.14661
zn          -2.473e-02  4.408e-03  -5.611  3.35e-08 ***
indus        5.722e-02  1.997e-02   2.865  0.00434 **
chas        -2.824e-01  2.846e-01  -0.992  0.32152
nox         -1.050e+01  1.187e+00  -8.848 < 2e-16 ***
rm          -7.076e-02  1.479e-01  -0.478  0.63255
age          7.198e-03  4.313e-03   1.669  0.09577 .
dis         -2.187e-02  6.883e-02  -0.318  0.75084
rad          1.177e-01  2.154e-02   5.465  7.35e-08 ***
tax          6.983e-04  1.244e-03   0.561  0.57491
black        1.573e-03  8.873e-04   1.773  0.07692 .
lstat       -3.770e-02  1.824e-02  -2.067  0.03929 *
medv        -1.021e-01  1.402e-02  -7.283  1.31e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.554 on 492 degrees of freedom
Multiple R-squared:  0.4982,    Adjusted R-squared:  0.485
F-statistic: 37.58 on 13 and 492 DF,  p-value: < 2.2e-16
```

We see there Ptratio has strong relation with Zn, indus, nox, rad and medv

While comparing with linear regression some results are matching and not all also the attributes which are matching are not in details when compared to rules as each attribute is not further classified in low, mid and high , I found association rule to be better while trying to find better results with proper attributes as they provide insights for lift, support and confidence and we can take better decision

**question 4 )**

**(10 points) (Modified Exercise 14.4) Cluster the demographic data of Table 14.1 using a classification tree. Specifically, generate a reference sample the same size as the training set. Build a classification tree to the training sample (class 1) and the reference sample (class 0) and describe the terminal nodes having highest estimated class 1 probability.**

Solution 1)

Step1 : Loading Marketing Data

```
library(ElmStatLearn)
library(rpart)
data("marketing")

## marketing data
new_data=marketing

str(new_data)

head(new_data)
```

Step 2 : Replacing 2694 NA values with median

```
> sum(is.na(market_data))
[1] 2694

> sum(is.na(market_data))
[1] 0
```

Step 3 : generating random value for all the variable using sample and adding target column with 1

Step 4 : Replicating the data and sampling all the variable and adding target column with value 0

step 5 : Merging the both the data frame using rbind

```
final_data = rbind(new_data_frame, data_frame);
```

step 6 : Converting all the categorical features as factor

Step 7 : Making model using rpart and seeing summary

```
> model = rpart(target~., final_data)
> summary(model)
Call:
rpart(formula = target ~ ., data = final_data)
n= 17986

   CP nsplit rel error xerror xstd
1 0.01      0      1      0      0

Node number 1: 17986 observations
  predicted class=0  expected loss=0.5  P(node) =1
  class counts: 8993 8993
  probabilities: 0.500 0.500

> |
```

Step 8: Predicting the values

```
> model.predict = predict(model, final_data[, -c(15)])
> model.predict
      0      1
[1,] 0.5 0.5
[2,] 0.5 0.5
[3,] 0.5 0.5
[4,] 0.5 0.5
[5,] 0.5 0.5
[6,] 0.5 0.5
[7,] 0.5 0.5
[8,] 0.5 0.5
[9,] 0.5 0.5
[10,] 0.5 0.5
[11,] 0.5 0.5
```

As we can see the node value is 0.5 we cannot reach to any conclusion and using rpart is not advisable for this dataset.